




The AI ethics trust engine

Converting capital into capabilities

In collaboration with the Notre Dame–IBM Tech Ethics Lab

How IBM can help

Clients can realize the potential of AI, analytics, and data using IBM's deep industry, functional, and technical expertise; enterprise-grade technology solutions; and science-based research innovations.


For more information:

AI services from IBM Consulting®
ibm.com/consulting/artificial-intelligence

AI solutions from IBM Software
ibm.com/watsonx

AI innovations from IBM Research®
research.ibm.com/artificial-intelligence

The Notre Dame–IBM Tech Ethics Lab
techethicslab.nd.edu



Organizations don't have to choose between doing AI right and doing the right AI.

Key takeaways

- **Organizations see measurable business results when AI ethics becomes integral to AI value creation.**

Investments in AI ethics correlate with increased operating profit from AI. Organizations in the top quartile of AI ethics spending, as a percent of AI spend, saw 34% higher operating profit attributable to AI in 2023 versus those in the lowest quartile. Executives also highlight the top three benefits from investments in AI ethics: increased trust (61%), strengthened brand reputations (57%), and mitigated reputational risks (54%).
- **Trust issues are stalling AI adoption.**

More than half of executives say ethics-oriented concerns are key barriers and building trust is a significant internal challenge. AI that is not trusted does not get used, undermining even the most robust business case.
- **AI ethics approaches often fail when treated as a single department's responsibility.**

53% of executives say their governing bodies for AI ethics are not effective. For many, accountability is concentrated in leadership roles, but it's important for multiple operational teams to influence AI ethics outcomes. Leading practices include AI ethics committees with cross-functional representation and tooling to coordinate inputs from key stakeholders.

Contents

Introduction	3
Part one: What is the business case for investing in AI ethics?.....	5
Part two: What is limiting AI's business value?.....	12
Part three: How should organizations operationalize AI ethics for optimal impact?	17
Action guide: Four starting points, one destination—pathways to approaching AI ethics holistically	22

Introduction

How trustworthy AI drives business performance

For years, enterprise leaders have debated whether investing in AI ethics is a business nicety or a business necessity. Some skeptics even raise concerns about it slowing adoption and time to value. New research from the IBM Institute for Business Value (IBM IBV) helps resolve this debate.

While some companies view responsible AI as a moral obligation and societal imperative regardless of other business impacts, many also report clear financial benefits as a consequence. In a new survey of 915 global executives—part of a three-year initiative sponsored by the IBM AI Ethics Board in collaboration with the Notre Dame–IBM Tech Ethics Lab and the IBM IBV—we found that organizations investing more in AI ethics consistently achieve higher operating profit, stronger return on investment (ROI), and measurable competitive advantage from AI.

This data validates an emerging finding highlighted in the first IBM IBV report contributing to this initiative, *Why invest in AI ethics and governance?* Based on interviews with over three dozen companies, the report suggests that organizations aren't choosing between business outcomes and doing the right thing. Those taking a holistic approach that combines practical business discipline with genuine ethical commitment are, in fact, asking a different question. It's not "should we invest in AI ethics?" It's "how do we integrate ethical practices into AI value creation?" They show that organizations embedding trust can capture more business value from their AI investments—including direct financial benefits, a positive brand and culture, and new capabilities that deliver long-term advantages.¹


Yet many companies remain trapped between good intentions and lagging execution, preventing them from capturing this value. Only about one-third of the executives in our new study say their organizations are applying a set of core AI ethics-specific tools today. And more than half cite trust, bias, and explainability as constraints to AI adoption. While recent data points to early signs that AI adoption is slowing,² and organizations are resetting value expectations,³ the costs of ignoring these barriers have already materialized—from a Berlin bank's €300,000 fine for denying a credit application via AI without GDPR-required transparency to public backlash when a gaming company's AI-based anti-cheat system erroneously banned legitimate players.⁴ Leaders cannot afford inaction—especially where business cases are strong. The promise of agentic AI—which also comes with greater autonomy and complexity—only amplifies the urgency to address these challenges strategically.

This report explores key questions every leader must answer to realize the value of trustworthy AI:

Part 1 What is the business case for investing in AI ethics?	Part 2 What is limiting AI's business value? Why do current approaches to trustworthy AI fall short?	Part 3 How should organizations operationalize AI ethics for optimal impact? How can accountability and collaboration make AI ethics a business accelerator?	Action Guide Pathways to approaching AI ethics holistically.
--	---	---	--

Definitions⁵

The following terms can be used in different ways, and while we're not necessarily prescribing the following definitions—or the set of principles companies adopt, because cultural and organizational norms vary⁶—we share them as a reference point.
AI ethics: A multidisciplinary field that studies how to optimize AI's beneficial impact while reducing risks and adverse outcomes. Examples of AI ethics issues are data responsibility and privacy, fairness, explainability, robustness, transparency, environmental sustainability, inclusion, moral agency, value alignment, accountability, trust, and technology misuse.
AI governance: An organization's act of governing, through its corporate instructions, staff, processes, and systems, to direct, evaluate, monitor, and take corrective action throughout the AI lifecycle, and to provide assurance that an AI system is operating as the organization intends, as its stakeholders expect, and as required by relevant regulation.
Trustworthy AI: AI systems that help to build trust through a human-centric approach characterized by five requirements: explainability, fairness, robustness, transparency, and privacy.
Responsible AI: Trustworthy AI that is used in an environmentally and socially responsible way.



Part one:

What is the business case for investing in AI ethics?

Those who invest more see more business value.

In our 2022 study, *AI Ethics in Action*, early indications pointed to the value of investments in AI ethics—specifically, those organizations that placed a greater importance on AI ethics reported a higher degree of trust from customers and employees.⁷ But isolating the signal from the noise proved challenging.


Now with our new research, the signs are much clearer: investing in AI ethics directly correlates with superior business performance (see Figure 1). Organizations in the top quartile of AI ethics spending consistently demonstrate higher operating profit attributable to AI—an average of 30% higher than the lowest quartile over the past two years—while investments across all quartiles correlate with increasing operating profit.

This trend also applies to other measures of financial value from AI, including ROI. While investments in AI ethics explain only a relatively modest portion of AI's benefit—a large amount would be surprising indeed—it is statistically significant when isolated, most likely linked to organizations' AI maturity, as we observed in our 2023 research.⁸

Over the past two years, organizations in the top quartile of AI ethics spending demonstrate an average of 30% higher operating profit attributable to AI than the lowest quartile.

FIGURE 1

Organizations spending more on AI ethics credit greater profits to AI.



Based on IBM IBV analysis.


Organizations appear to be recognizing this value advantage. Executives report they plan to double their spending on AI ethics as a percent of total AI spend from 2023 to 2026. What's more, organizations consistently exceed their own AI ethics spending predictions, a pattern we don't usually see in other technology investments (see Figure 2).⁹

In 2024 alone, actual AI ethics spending reached 7% of AI spend versus the projected 4.3%, suggesting that early results compel increases. With trust and ethical concerns as persistent barriers to AI adoption, this accelerated investment represents a strategic move to unlock broader AI value rather than simply mitigate risk.¹⁰

FIGURE 2


Organizations consistently revise AI ethics spending upward from original projections.

Percentages represent AI ethics spend as a percent of AI spend



Note: Organizations estimate AI spend is approximately 10%–12% of IT spend.

Source: Projected percentages are unpublished data from the IBM IBV 2024 report, The ingenuity of generative AI; actual spend and revised projection are from our 2025 survey.




From values to value across multiple dimensions

Investments in AI ethics deliver business value through three primary channels: economic returns, reputational impact, and future capabilities.¹¹ Most executives (59%) say their AI ethics efforts are showing results, whether that's stronger customer trust, fewer complaints, or higher-quality products. More specifically, the top benefits they cite paint a picture of fortifying the business end to end by creating market differentiation, protecting against costly failures, and building competitiveness (see Figure 3).

FIGURE 3

Trustworthy AI delivers end-to-end business value—from trust building to risk mitigation.

Percent saying organization has realized each benefit from investing in AI ethics




Q. What specific business benefits have you realized from your investment in AI ethics?

The performance improvements from these benefits are concrete (see Figure 4). And in terms of regulatory costs alone, organizations that say AI ethics investments have generated savings report significantly higher ROI, an advantage that is growing year-over-year.


FIGURE 4

AI ethics improves key performance metrics.

Percent improvement reported for the last 12 months



Q. For the last 12 months, how much have the following metrics improved due to AI ethics?



Organizations have different priorities for investing in AI ethics. Our analysis reveals four distinct strategic approaches. The largest group is brand centric, viewing AI ethics as a market differentiator that strengthens external stakeholder relationships and positions them as trusted partners. The second group emphasizes building trust to accelerate internal AI adoption. A third group is growth driven, pursuing AI ethics primarily for revenue gains. Notably, a smaller segment pursues value creation across all domains—compliance management, AI adoption, reputation, and financial performance. The holistic-focused approach generates the largest gains: this group projects the highest operating margin impact of the four groups by 2027 (17.2%), lower regulatory penalties, and the most improved resilience to regulatory changes. This finding validates our initiative's motivating hypothesis on the value of a holistic approach to understand the return on AI ethics investments,¹² refined and supported by subsequent research with more than three dozen organizations.¹³

Organizations prioritize AI ethics for different reasons. Most view it as a market differentiator.

Case study

Global financial services firm: Justifying positive returns with a lowered reputational risk¹⁴


For one senior leader responsible for AI governance at a global financial services firm, AI ethics practices start with education. He advocates hosting workshops that discuss ethical principles and values—empowering leadership to understand trade-offs. “We need to talk about AI in a way that interests leadership, not just in processes and procedures,” he observes.


In discussing how to measure the return on investments in AI ethics, the senior leader offers the “creepy line” metaphor. Often, organizations find themselves in situations in which they are doing something perfectly legal that is highly profitable yet still feel uncertain about the ethicality of their actions—a sense of crossing the “creepy line.”

In such situations, he says that organizations must examine the activity through the lens of both current and future generations, in conjunction with all comprehensive ethical considerations. As long as these considerations are covered satisfactorily, the organization should feel reassured that the “creepy line” is not breached.

He also notes, “Reputational risk is a key factor in justifying positive returns. We aim to decrease reputational risk while applying data and AI ethics principles.” For example, his team conducted an ethical fairness review of loan pricing involving a credit scoring algorithm. In conducting this review, the team analyzed all 165 features of the model, asking if there were any potential causal mechanisms for why that particular data feature may correlate with an individual’s ability to pay back a loan.

Ultimately, three data features were removed because a causal link did not exist, thus avoiding the lack of fairness in using this AI technology.





Part two:

What is limiting AI's business value?

Trust issues and modest implementation of ethics practices are stalling AI adoption.

The connection between AI ethics investment and business performance runs through a critical pathway: trust. After all, AI that is not trusted does not get used by employees, customers, partners, or other stakeholders—undermining even the most robust business case.

Previous IBM IBV research showed a lack of trust is limiting AI's adoption and potential.¹⁵ And our latest research confirms the problem persists. More than half (56%) of organizations cite explainability, ethics, bias, or trust as key barriers to AI adoption—meaning more than half of enterprises are leaving AI value on the table due to concerns related to its trustworthiness. The ascent of agentic AI only intensifies risks and uncertainty (see Perspective “Navigating the agentic AI risk landscape” on page 15). Nearly two-thirds (65%) of executives say agentic AI will require stricter ethical guidelines than current AI systems, and 64% say they will need to rethink their AI frameworks significantly as AI systems gain more autonomy. Until trust concerns are addressed, they will endure as fundamental business blockers, preventing organizations from scaling AI use cases that could drive competitive advantage, operational efficiency, and revenue growth. In fact, the question of operational efficiency triggered the IBM AI Ethics Board to launch this broader initiative to explore the link between AI ethics and business outcomes.

65% of executives say agentic AI will require stricter ethical guidelines than current AI systems.

Implementation reality check

Almost two-thirds (64%) of executives who view trust and ethical concerns as a major roadblock to adoption say that building trust is a significant internal issue. Over half of enterprises understand that AI ethics issues matter and impact the business, but most lack the operating models to address it effectively:

- Only **41%** have established approaches for integrating AI ethics into their AI strategy.
- **62%** report regular tension between business goals and ethical values.
- Only **49%** report fully transparent AI decision-making processes.

Diving deeper, implementation of AI ethics best practices is limited. While basic visible frameworks and some foundational tools may be in place, most organizations lack advanced operational capabilities such as impact assessments and ambassadors embedded throughout the business (see Figure 5). These tools translate principles into practice. Without them, organizations may operate in ethical blind spots, missing market opportunities, deploying AI without adequate governance, and exposing risks.

Policy limitations are then exacerbated by talent and funding challenges:

- **67%** struggle with availability of skilled personnel.
- **65%** find upskilling and training at scale challenging.
- **60%** cite lack of financial resources as a key obstacle.

Many organizations are trapped in a vicious cycle: they can't build trust in AI because they lack the capabilities to implement ethical practices, yet they struggle to justify investing in those capabilities because they haven't demonstrated their value. Without a new approach, costs will continue to compound, AI adoption will slow, and stakeholder confidence will ebb.


The path forward requires treating principles as starting points, not destinations—immediately pairing each ethical commitment with specific operational mechanisms that make it actionable. By building capabilities in parallel, organizations can compress the ethics implementation timeline and start capturing value.

Without a new approach to AI ethics, costs will continue to compound, AI adoption will slow, and stakeholder confidence will ebb.

FIGURE 5

Organizations lag in adoption of advanced operational AI ethics practices.

Percent reporting some or full adoption of each practice



Q. To what degree have you adopted the following AI ethics-specific mechanisms today?

Perspective

Navigating the agentic AI risk landscape¹⁶

Agentic AI—highly autonomous systems that complete tasks without continuous human oversight—introduces unprecedented risks that traditional AI governance frameworks weren’t designed to handle.

Unlike conventional AI systems, agentic AI exhibits four characteristics that amplify risk exposure—categories identified in research conducted by IBM’s AI Ethics Board:


- **Opaqueness.** Limited visibility into decision-making processes.
- **Open-endedness.** Self-selection of resources, tools, and even other AI agents.
- **Complexity.** Increasingly sophisticated learning and adaptation capabilities.
- **Non-reversibility.** Higher likelihood of taking irreversible actions with real-world consequences.

These characteristics create entirely new risk categories. For example, an AI agent might autonomously modify datasets in ways that introduce bias at scale, or it could inappropriately share confidential information across systems without human awareness.

As AI agents proliferate rapidly across enterprises, risks become more challenging. Centralized repositories that provide pre-vetted, governed agents can help. But mitigation requires an end-to-end governance approach that extends beyond traditional AI risk management. Key strategies include:

- **Human-in-the-loop** validation for critical decisions.
- **Restricted resource access** with defined boundaries.
- **Real-time monitoring** of agent actions and interactions.
- **Reversibility protocols** where technically feasible.

The challenge is that agentic AI is advancing faster than best practices can emerge. Organizations must adapt existing responsible AI frameworks while developing new governance capabilities specifically designed for autonomous systems.



Case study

SAS: Ensuring an AI-driven future that is built for all of us¹⁷

Reggie Townsend, VP of the Data Ethics Practice at SAS, leads a team tasked with coordinating responsible innovation principles, operational workflows, and governance structures across a global organization. It all began with questions and investigating.

Prompted by risks to vulnerable populations and the increasing sophistication of AI, Townsend and close colleagues began digging deeper into responsible AI and data ethics at SAS. They were empowered by SAS leadership to formalize the company's longtime commitment to responsible innovation. Consequently, SAS created the Data Ethics Practice (DEP).


With a philosophy of "ethical by design," the DEP guides the company's efforts to help employees and customers deploy data-driven systems that promote human well-being, agency, and fairness. This approach compels individuals to answer three basic questions:


- For what purpose?
- To what end?
- For whom might it fail?

The team helps build trustworthy AI capabilities and workflows in support of customers' and developers' pursuit of their responsible AI goals. AI governance advisory services from the DEP are helping customers put AI into action responsibly. The DEP also provides critical counsel to employees on product development, marketing, and more.

When Townsend's role and team were created, the hope was their work would bolster trustworthiness of products, processes, and people. This, in turn, would enhance the brand's reputation as a trusted AI leader.

Profits are important, of course. But according to Townsend, his team's guiding principle is that wherever SAS software shows up, it does no harm. "Sometimes," Townsend observes, "you just have to take action because it's the right action to take."





Part three:

How should organizations operationalize AI ethics for optimal impact?

Enterprise-wide collaboration is essential for AI ethics to deliver results.

AI governance operates on two distinct levels: within executive leadership and across operational roles. To be effective, AI ethics needs to bridge that boundary. Yet too often, efforts are siloed.

Our research shows that primary accountability for AI ethics is concentrated across leadership functions, with no single entity dominating:

- AI ethics boards lead accountability in **19%** of organizations.
- Boards of directors hold responsibility in **16%** of cases.
- Risk and compliance functions own it in **13%** of organizations.


This data exposes a potential structural flaw for some organizations. While accountability is concentrated in leadership roles, the real power to shape AI ethics outcomes is distributed across multiple operational teams—and with diffuse reporting relationships (see Figure 6):

- Data teams control the foundation of ethical AI—the inputs that determine everything downstream.
- Risk and compliance create operational safeguards.
- Research, product design, and technology development embed ethics into solution architecture.
- Data science/machine learning/AI teams implement ethical considerations in model development.
- Information security protects against failures and breaches through technical controls.


FIGURE 6

AI ethics requires cross-functional collaboration.

Key functions shaping AI ethics: percent reporting significant influence or leadership roles



Q. To what extent do the following functions influence AI ethics at your organization?




The risk of misalignment suggests why 53% of organizations find their AI ethics governing bodies ineffective—they're trying to govern a cross-functional capability through single-function accountability models.

If AI governance is a siloed effort, ethics guidance can be detached from operational reality—creating friction rather than value. Conversely, when board-level oversight combines with cross-functional collaboration, ethical considerations become embedded in AI decisions rather than applied as external constraints.

This integration enables:

- Faster deployment because ethical considerations are built in, not bolted on.
- Better risk management because multiple functions contribute specialized expertise.
- More innovation because trustworthy AI constraints become design parameters that drive creative solutions.
- Stronger stakeholder trust because trustworthy AI becomes demonstrably operational, not just aspirational.

53% of organizations find their AI ethics governing bodies ineffective—they're trying to govern a cross-functional capability through single-function accountability models.



Building the collaborative framework


The more effective organizations create structured collaboration between governance and operations functions. They establish clear roles where boards and ethics committees set strategic direction and resource priorities, while operational functions translate those priorities into daily practices—with clear feedback loops, especially because direct reporting relationships are difficult in cross-functional organizations.

This structure requires organizations to create a culture that views trustworthy AI as spanning the entire AI lifecycle. For example, data teams must understand ethical implications of their choices. Product teams must balance user experience with ethical considerations. Risk teams must enable rather than constrain innovation.

When executed well, linking ethics into your governance structure accelerates innovation as it reduces uncertainty by increasing predictability. Teams often avoid deploying AI in higher-risk scenarios—such as using AI agents for financial reporting—because they lack clear guidance on safe deployment practices. Our previous research confirms the point: 75% of executives said the absence of a comprehensive external regulatory framework is a key barrier to AI adoption, so they look internally.¹⁸ Effective AI governance that incorporates ethical practices provides the “safe space” and clear guardrails that enable rapid scaling. Recognizing that not all AI applications carry equal risk, smart governance can take a tiered approach—streamlined processes for lower-risk use cases such as content generation or data analysis, and intensive oversight for high-risk applications involving financial decisions, healthcare, or customer-facing autonomous systems.

No obvious best practice for which groups should be accountable and responsible exists—so enterprises should focus on the best fit for their organizations. For example, recent IBM IBV research shows that those companies that have appointed a Chief AI Officer seem to nest both accountability and responsibility for AI ethics in that role.¹⁹

The organizations that will win with trustworthy AI know that collaboration transforms ethical intentions into business results. They’ve moved beyond asking whether they should invest in AI ethics to optimizing how they organize that investment for greatest impact. Shared values that benefit both society and business result from adopting a holistic mindset toward investment in AI ethics and governance.²⁰



Case study

Fidelity Investments: Making AI ethics a foundational pillar²¹


Responsible AI initiatives are embedded into each phase of AI use cases at Fidelity Investments, beginning with robust data management practices and feeding into a dynamic review process driven by the company's AI Center of Excellence. The financial services firm invested heavily in these initiatives to make AI ethics one of its foundational pillars—rather than a compliance box-checking exercise.

Each business line at Fidelity has a dedicated team for AI use case development and vendor management. This work is guided by the expertise of external consultants and actively monitored by the firm's compliance and risk officers, who receive specialized AI training.

The AI Center of Excellence is involved in each step of this process, from vendor selection to model evaluation. It resides in Fidelity's data function and includes representation from each business unit at the firm, with roles ranging from risk compliance and audit to legal and even information security. This process also allows Fidelity to confidently answer clients' increasing demands for information on its AI use and governance.

Resistance to responsible AI initiatives is inevitable, as they can delay projects or limit use cases. "You have to explain that the reason controls are so important is not just some random compliance policy, but that there are implications to the firm if we get this wrong," says a senior vice president with responsibility for data policy at the company. Fidelity has been able to minimize pushback by framing these initiatives as integral to the success of AI projects and by streamlining the overall governance process.

Fidelity's AI Center of Excellence resides in the data function and includes representation from each business unit at the firm.



Action guide

Four starting points, one destination—pathways to approaching AI ethics holistically


Our ongoing adjacent research suggests that organizations typically justify AI ethics investments based on their primary stakeholder priorities. The following actions are tailored to the four distinct motivations for investing in AI ethics identified in this research. They are designed to help you advance both your AI and AI ethics maturity while aligning with your strategic objectives and organizational focus.




Brand-differentiating

Organizations investing in AI ethics to proactively bolster their brand differentiation and overall reputation

Your external focus helps you stand out in the market—concentrate on those values important to driving your differentiation, then amplify it through transparent communication about your tangible AI practices. Conduct customer/employee research on which of your values set you apart, align your AI ethics competencies to your AI strategy, and then publish regular transparency reports that detail your concrete approaches, challenges overcome, and lessons learned, positioning your organization as a thought leader. Enhance brand credibility by incorporating customer-facing explainability, privacy, and personalization features that translate organizational investments into clear product value. Measure results on the defined metrics. Partner with industry associations and academic institutions to co-develop standards, collaboratively aiming to advance the field of AI ethics in your industry with research, open publications, toolkits, use cases, working groups, and so on—enhancing your leadership position while shaping policies and practices that benefit the greater societal good. Be mindful of the “ethics washing” trap; don’t just share practices externally—actually deliver consumer and employee benefits by differentiating proactively and meaningfully. Then consider what steps can be taken to augment your current strengths by adjusting your orientation toward growth and AI adoption.






Adoption-oriented

Organizations investing in AI ethics to build internal trust and facilitate broader adoption of AI


Your success at building internal trust positions you to capture broader business value. Continue to build on a robust foundation of employee AI literacy through advanced training, distinct human-in-the-loop approaches, and ongoing tailored role- and function-based education. Ensure transparent and user-friendly processes for ethics approvals, including an open backlog and open descriptions on why cases are approved or not, and what mitigation measures are taken so teams can learn for the next AI project. Test human-in-the-loop interactions and enhance where needed, as the human-AI team might perform differently than you expect in practice.²² Invest in making these processes easy to use, smooth, partially automated, and transparent—with positive reinforcement for those who highlight ethics issues—empowering employees. Most importantly, leverage employee confidence in your AI systems to drive innovation. Encourage teams to propose bold AI applications knowing they have the ethical infrastructure to deploy them safely and quickly. Then take action to develop capabilities focused on strengthening your brand and fueling growth to balance your concentration on adoption.



Growth-oriented

Organizations investing in AI—and AI ethics—to fuel primarily revenue growth


Your revenue-driven approach to AI ethics creates a powerful competitive advantage. Translate your internal AI wins into customer-facing value propositions, showing how your trustworthy AI practices deliver better products and services. Accelerate time-to-market for new AI capabilities, turning ethical deployment processes into a speed differentiator. Convert AI ethics competencies into product, services, or solution advantages that deliver sustainable growth. For example, develop customer-facing features, such as bias transparency dashboards or explainability interfaces, and craft the metrics to measure their success qualitatively and quantitatively. At the same time, sharpen your focus on compliance—regulatory violations can quickly undermine external trust and erode the revenue gains you’re capturing. Invest in baseline compliance infrastructure to protect your investments. Drive exponential growth by expanding into markets where compliance is becoming a procurement requirement, turning regulatory readiness into a competitive edge. Then cultivate capabilities directed toward fortifying your brand and adoption for a more holistic approach.



Holistic-focused

Organizations investing across the board in AI ethics to manage regulatory compliance, boost their brand, fuel growth, and build internal trust and adoption

Your holistic approach positions you for maximum impact—optimize it through integrated measurement and cross-functional alignment. Develop unified dashboards that track ethics outcomes across all dimensions (compliance, adoption, reputation, financial performance) to identify which investments deliver the highest compound returns. Create cross-functional project teams where compliance, product, marketing, and technical teams collaborate on ethics initiatives from inception, ensuring every investment delivers value across multiple objectives. Use your comprehensive perspective to pilot innovative approaches such as trustworthy AI product features that can be monetized directly.



About the authors

Nicholas Berente

Senior Associate Dean for Academic Programs
Professor of IT, Analytics, and Operations
University of Notre Dame,
Mendoza College of Business
[linkedin.com/in/berente/](https://www.linkedin.com/in/berente/)
nberente@nd.edu

Nicholas is the James H. Sweeny III and Alicia Sweeny Collegiate Professor of IT, Analytics, and Operations at the University of Notre Dame's Mendoza College of Business, where he serves as the Senior Associate Dean for Academic Programs. His research focuses on the organizational, institutional, and ethical implications of AI.

Mariadena Bevilacqua

PhD Student in Analytics
University of Notre Dame,
Mendoza College of Business
[linkedin.com/in/mariadena-bevilacqua-6848b9132/](https://www.linkedin.com/in/mariadena-bevilacqua-6848b9132/)
mbevilac@nd.edu

Mariadena is an Analytics PhD candidate at the University of Notre Dame's Mendoza College of Business. Her research examines how AI and other emerging technologies influence human behavior and societal outcomes in ethically complex contexts, drawing primarily on experimental methods.

Marianna Ganapini

Associate Professor, School of Data Science
and Philosophy
University of North Carolina at Charlotte
[linkedin.com/in/marianna-b-ganapini-769624116/](https://www.linkedin.com/in/marianna-b-ganapini-769624116/)
bergamam@union.edu

In addition to being an Associate Professor, Marianna is Co-Founder and CEO of LogicaNow, a consultancy specializing in responsible AI and AI governance. Her research and advisory work focus on trust, autonomy, and risk governance in AI, drawing on both philosophical inquiry and industry practice. Marianna is Faculty Director at the Montreal AI Ethics Institute, and a member of ISO and CEN-CENELEC JTC 21 working on AI standards. She has collaborated with leading universities and companies and is widely published in AI ethics and cognitive science.

Brian Goehring


Associate Partner, AI Research Lead
IBM Institute for Business Value
[linkedin.com/in/bcgoehring](https://www.linkedin.com/in/bcgoehring)
goehring@us.ibm.com

Brian brings over 25 years' experience in strategy consulting with senior-level clients across most industries and functions, collaborating with academics, clients, and other experts to develop data-driven insights and thought leadership. He received an AB in Philosophy from Princeton University with certificates in Cognitive Studies and German.

Francesca Rossi

IBM Fellow and AI Ethics Global Leader
IBM Research
[linkedin.com/in/francesca-rossi-34b8b95/](https://www.linkedin.com/in/francesca-rossi-34b8b95/)
Francesca.Rossi2@ibm.com

Francesca is an IBM Fellow at the T.J. Watson Research Center and is the IBM Global Leader for Responsible AI and AI Governance. In this role, she leads research projects to advance AI capabilities and co-chairs the IBM Responsible Technology Board. She is the past president of the Association for the Advancement of AI (AAAI) and a former board member of the Partnership on AI. Before joining IBM, she was a professor of computer science at the University of Padova, Italy.



Contributors

Phaedra Boinodiris

Global Leader for Trustworthy AI
IBM Consulting

Rachna Handa

Research Manager
IBM Institute for Business Value

Daria Simons

Managing Consultant—Trustworthy AI
IBM Consulting

From the IBM Institute for Business Value analytics and creative team

Sara Aboulhosn, Associate Creative Director

Benjamin Hurte, Design Director

Heba Nashaat, Data and Content
Management Manager

Tihomir Trifonov, Visual Designer


Joanna Wilkins, Editorial Lead

Acknowledgements

The authors and contributors would like to thank the IBM AI Ethics Board for sponsoring the initiative that prompted this study and related research across 2023-2025, as well as the Notre Dame–IBM Tech Ethics Lab (and Lab Administration Sara Berger, Heather Domin, Michael Infantine, Adam Kronk, Jungmin Lee, Megan McDermott, Nuno Moniz, Cathy Quinlan, as well as Steering Committee members Nick Fehring, David Go, Betsy Greytok, Christina Montgomery, Jeffrey Rhoads, Meghan Sullivan, and Jeffrey Welsner) for partially funding this initiative during that period.

The right partner for a changing world

At IBM, we collaborate with our clients, bringing together business insight, advanced research, and technology to give them a distinct advantage in today's rapidly changing environment.



Subscribe to our IdeaWatch newsletter

Just the insights. At your fingertips. Delivered monthly.

Brought to you by the IBM Institute for Business Value, ranked #1 in thought leadership quality by Source Global Research for the second consecutive year.

Research-based thought leadership insights, data, and analysis to help you make smarter business decisions and more informed technology investments.

Subscribe now: ibm.co/ideawatch



Related reports

Why invest in AI ethics and governance?

Five real-world origin stories.

IBM Institute for Business Value in collaboration with the Notre Dame–IBM Tech Ethics Lab. December 2024.
<https://ibm.co/roi-ai-ethics>

From AI projects to profits: How agentic AI can sustain financial returns.

IBM Institute for Business Value. June 2025.
<https://ibm.co/agentic-ai-profits>

AI ethics in action: An enterprise guide to progressing trustworthy AI.

IBM Institute for Business Value. March 2022.
<https://ibm.co/ai-ethics-action>

915


Executives

18

Industries

19

Countries



Study approach and methodology

The study employed a quantitative research approach to examine the integration and management of AI ethics within organizational contexts. The data collection was conducted through a comprehensive global survey administered in collaboration with Oxford Economics, gathering responses from 915 executives across 18 industries, 19 geographic locations (see map), and over 19 distinct job functions. The research pursued three primary objectives: to establish benchmark practices in AI ethics implementation, to investigate the underlying motivations—examining both external compliance drivers and internal value-based imperatives, and to quantitatively assess the perceived business impact and financial returns of AI ethics initiatives, thereby evaluating their tangible organizational value beyond theoretical propositions.

To provide an initial assessment of the relationship between AI ethics investment and financial performance, a descriptive analysis was conducted. Organizations were segmented into quartiles (Q1–Q4) based on their spending on AI ethics as a percentage of total AI spend—where spending encompassed training and education, tools and technology, personnel, external consulting, and research. To ensure the validity of this quartile-based segmentation, the underlying distribution of spending was tested for bimodality, confirming its suitability. For each quartile, the mean percentage change in business unit operating profit attributable to AI was calculated for the years 2023 (actual), 2024 (actual), 2025 (estimated), and 2026 (estimated).

A multistage analytical strategy was then employed to uncover key relationships within the dataset. The investigation began with foundational techniques, including correlation and factor analysis, to identify underlying patterns. A primary component of the

analysis was a K-means cluster analysis, chosen for its computational efficiency with large samples and its ability to produce clearly separated clusters. This analysis segmented organizations based on four standardized business justification dimensions: compliance, brand reputation, trust, and growth. Standardization was a critical step to prevent variables with larger original scales from dominating the results. This process yielded a four-segment solution that offered an optimal balance between statistical fit and conceptual interpretability.

To quantify the financial impact of AI ethics initiatives, multiple linear regression analysis was conducted. This analysis tested the relationship between reductions in regulatory costs achieved through AI ethics programs and the overall perceived return on investment from AI initiatives, providing empirical evidence for their business value.

IBM Institute for Business Value

For two decades, the IBM Institute for Business Value has served as the thought leadership think tank for IBM. What inspires us is producing research-backed, technology-informed strategic insights that help leaders make smarter business decisions. From our unique position at the intersection of business, technology, and society, we survey, interview, and engage with thousands of executives, consumers, and experts each year, synthesizing their perspectives into credible, inspiring, and actionable insights. To stay connected and informed, sign up to receive IBM IBV's email newsletter at ibm.com/ibv. You can also find us on LinkedIn at <https://ibm.co/ibv-linkedin>.

Notes and sources

- 1 Berente, Nicholas, Marialena Bevilacqua, Heather Domin, Marianna Ganapini, Brian Goehring, and Francesca Rossi. *Why invest in AI ethics and governance? Five real-world origin stories*. IBM Institute for Business Value in collaboration with the Notre Dame–IBM Tech Ethics Lab. December 16, 2024. <https://ibm.co/roi-ai-ethics>
- 2 Kahn, Jeremy. “Is the AI bubble about to burst? There’s mounting data suggesting it might be.” *Fortune*. September 9, 2025. <https://fortune.com/2025/09/09/is-the-ai-bubble-about-to-burst-theres-mounting-data-suggesting-it-might-be/>
- 3 Blair, Kate PhD, Francesco Brenna, Nicholas Fuller PhD, Brian Goehring, and Matt Sanchez. *From AI projects to profits: How agentic AI can sustain financial returns*. IBM Institute for Business Value. June 9, 2025. <https://ibm.co/agentic-ai-profits>
- 4 “Berlin SA imposed 300 000 euro fine against bank after automated rejection of credit card application.” European Data Protection Board. May 31, 2023. https://www.edpb.europa.eu/news/national-news/2023/berlin-sa-imposes-300-000-euro-fine-against-bank-after-lack-transparency_en; Litchfield, Ted. “It looks like Call of Duty’s anti-cheat is permabanning innocent players.” *PC Gamer*. December 15, 2022. <https://www.pcgamer.com/it-looks-like-call-of-dutys-anti-cheat-is-permabanning-innocent-players/>
- 5 “Glossary: Privacy, AI Ethics, and Data Governance terms and definitions.” IBM Office of Privacy and Responsible Tech internal website. Accessed September 12, 2025.
- 6 Clancy, R.F., Q. Zhu, and S. Majumdar. “Exploring AI ethics in global contexts: a culturally responsive, psychologically realist approach.” *AI Ethics*. September 16, 2025. <https://doi.org/10.1007/s43681-025-00821-6>
- 7 Goehring, Brian, Francesca Rossi, and Beth Rudden. *AI ethics in action: An enterprise guide to progressing trustworthy AI*. IBM Institute for Business Value. March 31, 2022. <https://ibm.co/ai-ethics-in-action>
- 8 Ashoori, Maryam, Brian Goehring, Timothy Humphrey, Mahmoud Naghshineh, and Cathy Rodenback Reese. *Generating ROI with AI: Six capabilities that drive world-class results*. IBM Institute for Business Value. May 9, 2023. <https://ibm.co/ai-capabilities>
- 9 Goehring, Brian, Manish Goyal, Ritika Gunnar, Anthony Marshall, and Aya Soffer. *The ingenuity of generative AI: Unlock productivity and innovation at scale*. IBM Institute for Business Value. June 11, 2024. <https://ibm.co/scale-generative-ai>. Unpublished data. Blair, Kate PhD, Francesco Brenna, Nicholas Fuller PhD, Brian Goehring, and Matt Sanchez. *From AI projects to profits: How agentic AI can sustain financial returns*. IBM Institute for Business Value. June 9, 2025. <https://ibm.co/agentic-ai-profits>. Unpublished data.
- 10 Khandabattu, Haritha. “The 2025 Hype Cycle for Artificial Intelligence Goes Beyond GenAI.” Gartner. July 8, 2025. <https://www.gartner.com/en/articles/hype-cycle-for-artificial-intelligence>
- 11 Berente, Nicholas, Marialena Bevilacqua, Heather Domin, Marianna Ganapini, Brian Goehring, and Francesca Rossi. *Why invest in AI ethics and governance? Five real-world origin stories*. IBM Institute for Business Value in collaboration with the Notre Dame–IBM Tech Ethics Lab. December 16, 2024. <https://ibm.co/roi-ai-ethics>
- 12 Bevilacqua, Marialena, Nicholas Berente, Heather Domin, Brian Goehring, and Francesca Rossi. “The Return on Investment in AI Ethics: A Holistic Framework.” *Proceedings of the 57th Annual HICSS Conference on Systems Sciences*. January 2024. <https://arxiv.org/abs/2309.13057>.
- 13 Domin, Heather, Francesca Rossi, Brian Goehring, Marianna Ganapini, Nicholas Berente, and Marialena Bevilacqua. “On the ROI of AI Ethics and Governance Investments: From Loss Aversion to Value Generation.” *California Management Review*. July 29, 2024. <https://cmr.berkeley.edu/2024/07/on-the-roi-of-ai-ethics-and-governance-investments-from-loss-aversion-to-value-generation/>. Ganapini, Marianna Bergamaschi, Nicholas Berente, Francesca Rossi, Brian Goehring, and Marialena Bevilacqua. “Making Sense of AI Ethics and Governance Investments.” Paper presented at the Eighth AAAI/ACM Conference on AI, Ethics, and Society, Madrid, Spain, October 20–22, 2025.
- 14 Berente, Nicholas, Marialena Bevilacqua, Heather Domin, Marianna Ganapini, Brian Goehring, and Francesca Rossi. *Why invest in AI ethics and governance? Five real-world origin stories*. IBM Institute for Business Value in collaboration with the Notre Dame–IBM Tech Ethics Lab. December 16, 2024. <https://ibm.co/roi-ai-ethics>
- 15 Goehring, Brian, Francesca Rossi, and Beth Rudden. *AI ethics in action: An enterprise guide to progressing trustworthy AI*. IBM Institute for Business Value. March 31, 2022. <https://ibm.co/ai-ethics-in-action>
- 16 *AI agents: Opportunities, risks, and mitigations*. IBM AI Ethics Board. March 2025. <https://www.ibm.com/downloads/documents/us-en/1227c12efb38b2b3>; IBM AI Ethics Board. “Getting your organization ready to scale responsible agentic AI.” IBM Think blog. April 11, 2025. <https://www.ibm.com/think/insights/scale-responsible-agentic-ai>
- 17 Berente, Nicholas, Marialena Bevilacqua, Heather Domin, Marianna Ganapini, Brian Goehring, and Francesca Rossi. *Why invest in AI ethics and governance? Five real-world origin stories*. IBM Institute for Business Value in collaboration with the Notre Dame–IBM Tech Ethics Lab. December 16, 2024. <https://ibm.co/roi-ai-ethics>

- 18 Blair, Kate PhD, Francesco Brenna, Nicholas Fuller PhD, Brian Goehring, and Matt Sanchez. *From AI projects to profits: How agentic AI can sustain financial returns*. IBM Institute for Business Value. June 9, 2025. <https://ibm.co/agent-ai-profits>. Unpublished data.
- 19 Al Falasi, Saeed, Lula Mohanty, Irfan Verjee, Anthony Marshall, and Jacob Dencik PhD. *2025 Global C-Suite Series. Solving the AI ROI puzzle. How Chief AI Officers cut through complexity to create new paths to value*. IBM Institute for Business Value in partnership with the Dubai Future Foundation. July 13, 2025. <https://ibm.biz/chief-ai-officer>
- 20 Ganapini, Marianna Bergamaschi, Nicholas Berente, Francesca Rossi, Brian Goehring, and Marialena Bevilacqua. “Making Sense of AI Ethics and Governance Investments.” Paper presented at the Eighth AAAI/ACM Conference on AI, Ethics, and Society, Madrid, Spain, October 20–22, 2025.
- 21 Berente, Nicholas, Marialena Bevilacqua, Heather Domin, Marianna Ganapini, Brian Goehring, and Francesca Rossi. *Why invest in AI ethics and governance? Five real-world origin stories*. IBM Institute for Business Value in collaboration with the Notre Dame–IBM Tech Ethics Lab. December 16, 2024. <https://ibm.co/roi-ai-ethics>
- 22 Gaudeul, Alexia, Ottla Arrigoni, Vasiliki Charisi, Marina Escobar Planas, and Isabelle Hupont Torres. *The Impact of Human-AI Interaction on Discrimination*. Publications Office of the European Union, Luxembourg, 2025. <https://publications.jrc.ec.europa.eu/repository/handle/JRC139127>

© Copyright IBM Corporation 2025

IBM Corporation
New Orchard Road
Armonk, NY 10504

Produced in the United States of America | October 2025

IBM, the IBM logo, ibm.com, IBM Consulting, and IBM Research are trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the web at “Copyright and trademark information” at: ibm.com/legal/copytrade.shtml.

This document is current as of the initial date of publication and may be changed by IBM at any time. Not all offerings are available in every country in which IBM operates.

THE INFORMATION IN THIS DOCUMENT IS PROVIDED “AS IS” WITHOUT ANY WARRANTY, EXPRESS OR IMPLIED, INCLUDING WITHOUT ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND ANY WARRANTY OR CONDITION OF NON-INFRINGEMENT. IBM products are warranted according to the terms and conditions of the agreements under which they are provided.

This report is intended for general guidance only. It is not intended to be a substitute for detailed research or the exercise of professional judgment. IBM shall not be responsible for any loss whatsoever sustained by any organization or person who relies on this publication.

Examples presented are illustrative only. Actual results will vary based on client configurations and conditions and, therefore, generally expected results cannot be provided.



<https://ibm.biz/ai-ethics-bc>

